



Per-axis model projections (mean $\pm$ 95% bootstrap CI over prompts)


Refusal Stiffness




Sycophancy




Hedge Density




Harmfulness




Warmth



Bluntness



Verbosity



How to read

Each panel shows one persona axis. Dots mark each model's mean projection over all evaluation prompts; the line connects to zero (the cross-prompt midpoint).

- Horizontal bars are 95% bootstrap confidence intervals (resampled over prompts).
- Models are grouped by family (color); families separated by a thin gap.

Look for: (a) which models are consistently far from zero on an axis, (b) whether models within a family cluster together, (c) axes where CIs overlap heavily — those don't discriminate models.