



Per-layer persona-axis projections


Refusal Stiffness




Sycophancy




Hedge Density




Harmfulness




Warmth



Bluntness



Verbosity



How to read

Each panel shows how strongly each model's responses project onto a persona axis as we move through the prober's layers (Gemma-4).

- Curves trace mean projection across prompts at each layer.
- If lines stack on top of each other, it means the prober's per-layer geometry dominates over model-to-model differences.

See the residual variant (04b) for the model-specific signal alone.

command-a
tiny-aya-global

deepseek-r1-distill-qwen-7b
gemini-3.0-flash

Model
gpt-5.4
gpt-oss-20b

llama-3.1-8b-instruct
mistral-7b-instruct-v0.3

qwen-2.5-7b-instruct