



Per-layer model-specific signal (residual after subtracting cross-model mean)

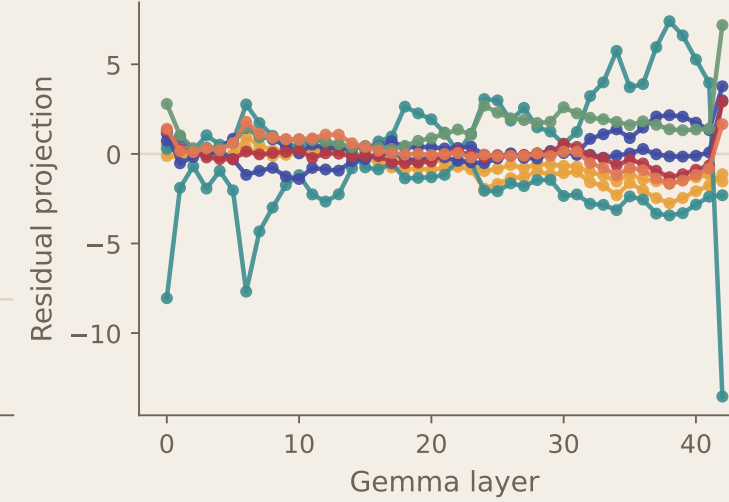
Refusal Stiffness



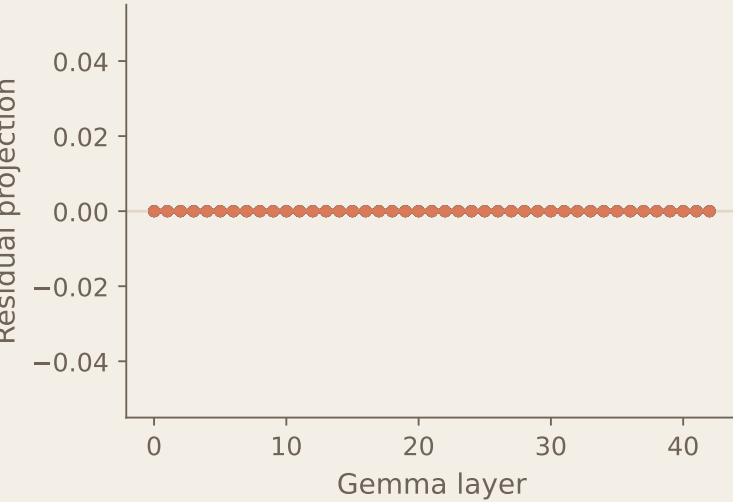
Sycophancy



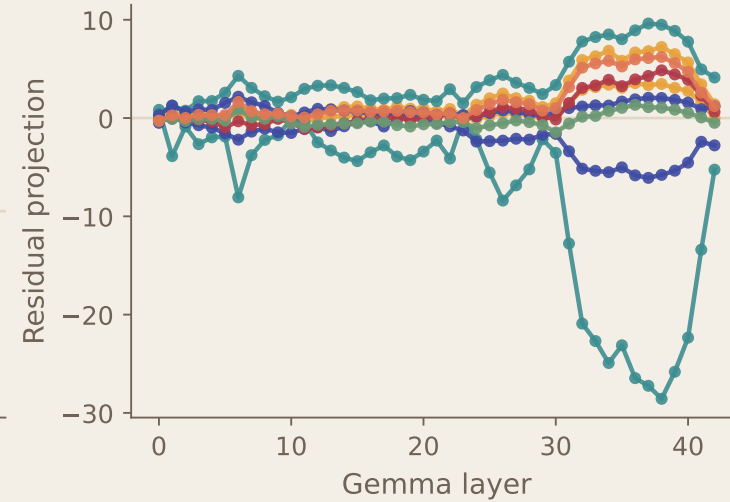
Hedge Density



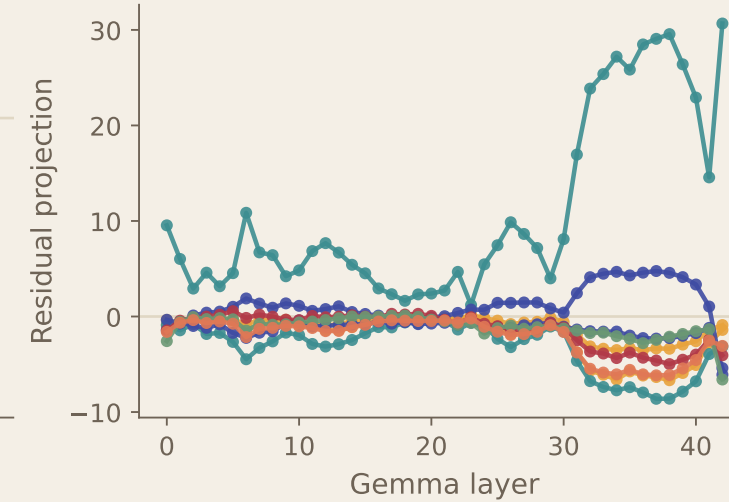
Harmfulness



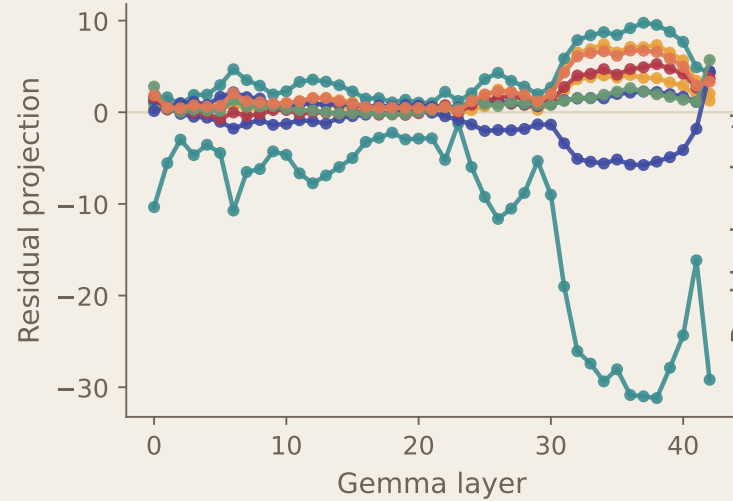
Warmth



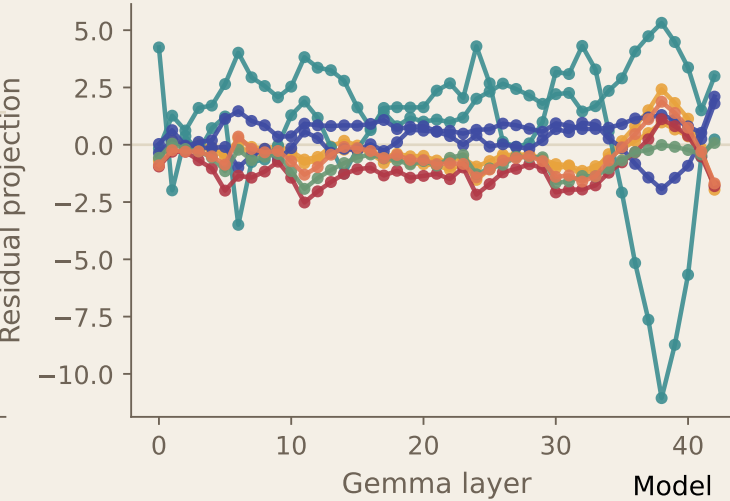
Bluntness



Verbosity



Creature Whimsy



How to read

Same as the raw per-layer plot, but with the cross-model mean subtracted at every (axis, layer). What remains is the model-specific signal.

- Lines above zero → the model is above peers at that layer for that axis.
- Below zero → below peers.

Look for layers where models fan out (high discriminability) — those are the layers carrying family-level signal. Flat-near-zero layers carry no information.

