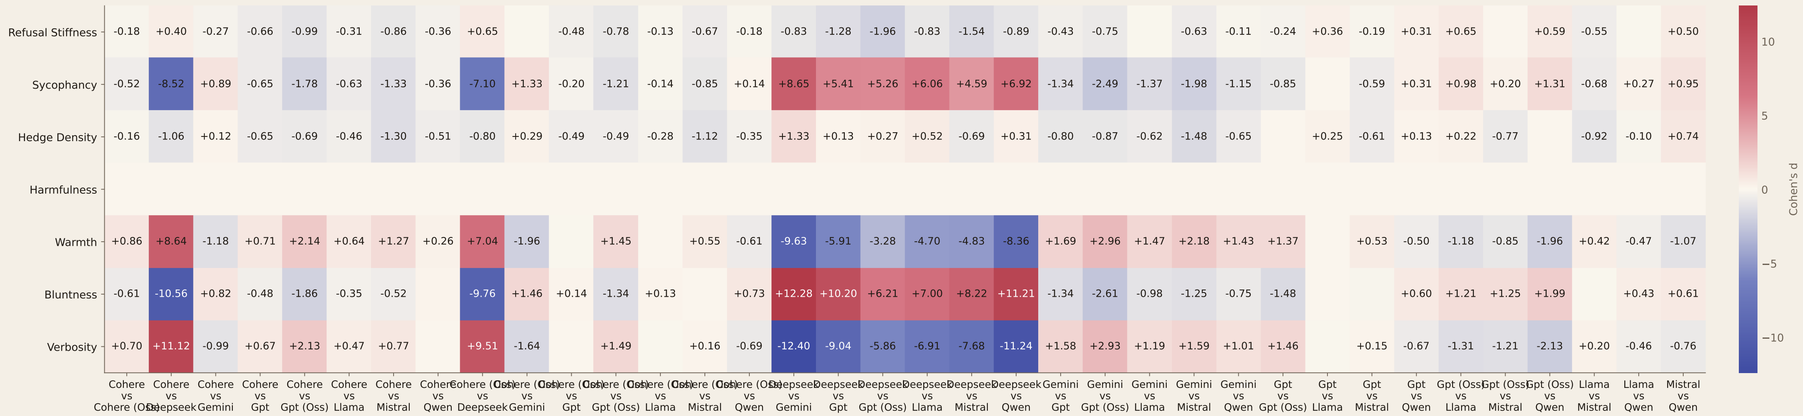

How to read

Each cell is the standardized mean difference (Cohen's d) between two families on one persona axis, computed over per-prompt projection scores.

- Red → row-family scores higher than column-family on that axis.
- Blue → row-family scores lower.
- |d| ≈ 0.2 small, 0.5 medium, 0.8 large.

Reads as: 'on axis X, family A is +d standard deviations higher than family B'. Big colored cells highlight the actual family-discriminating axes.